

Quicksort on Arrays

Quicksort shines for sorting arrays. In-place quicksort is very fast. But a fast in-place quicksort is tricky to code. It's easy to write a buggy or quadratic version by mistake. Goodrich and Tamassia did.

Suppose we have an array *a* in which we want to sort the items starting at *a[low]* and ending at *a[high]*. We choose a pivot *v* and move it out of the way by swapping it with the last item, *a[high]*.

We employ two array indices, *i* and *j*. *i* is initially "low - 1", and *j* is initially "high", so that *i* and *j* sandwich the items to be sorted (not including the pivot). We will enforce the following invariants.

- All items at or left of index *i* have a key \leq the pivot's key.
- All items at or right of index *j* have a key \geq the pivot's key.

To partition the array, we advance the index *i* until it encounters an item whose key is greater than or equal to the pivot's key; then we decrement the index *j* until it encounters an item whose key is less than or equal to the pivot's key. Then, we swap the items at *i* and *j*. We repeat this sequence until the indices *i* and *j* meet in the middle. Then, we move the pivot back into the middle (by swapping it with the item at index *i*).

An example is given at right. The randomly selected pivot, whose key is 5, is moved to the end of the array by swapping it with the last item. The indices *i* and *j* are created. *i* advances until it reaches an item whose key is \geq 5, and *j* retreats until it reaches an item whose key is \leq 5. The two items are swapped, and *i* advances and *j* retreats again. After the second advance/retreat, *i* and *j* have crossed paths, so we do not swap their items. Instead, we swap the pivot with the item at index *i*, putting it between the lists *I1* and *I2* where it belongs.

What about items having the same key as the pivot? Handling these is particularly tricky. We'd like to put them on a separate list (as we did for linked lists), but doing that in place is too complicated. As I noted previously, if we put all these items into the list *I1*, we'll have quadratic running time when all the keys in the array are equal, so we don't want to do that either.

The solution is to make sure each index, *i* and *j*, stops whenever it reaches a key equal to the pivot. Every key equal to the pivot (except perhaps one, if we end with *i* = *j*) takes part in one swap. Swapping an item equal to the pivot may seem unnecessary, but it has an excellent side effect: if all the items in the array have the same key, half these items will go into *I1*, and half into *I2*, giving us a well-balanced recursion tree. (To see why, try running the pseudocode below on paper with an array of equal keys.) WARNING: The code on page 530 of Goodrich and Tamassia gets this WRONG. Their implementation has quadratic running time when all the keys are equal.

```
public static void quicksort(Comparable[] a, int low, int high) {
    // If there's fewer than two items, do nothing.
    if (low < high) {
        int pivotIndex = random number from low to high;
        Comparable pivot = a[pivotIndex];
        a[pivotIndex] = a[high];           // Swap pivot with last item
        a[high] = pivot;

        int i = low - 1;
        int j = high;
        do {
            do { i++; } while (a[i].compareTo(pivot) < 0);
            do { j--; } while ((a[j].compareTo(pivot) > 0) && (j > low));
            if (i < j) {
                swap a[i] and a[j];
            }
        } while (i < j);

        a[high] = a[i];
        a[i] = pivot;                       // Put pivot in the middle where it belongs
        quicksort(a, low, i - 1);           // Recursively sort left list
        quicksort(a, i + 1, high);         // Recursively sort right list
    }
}
```

Can the "do { i++ }" loop walk off the end of the array and generate an out-of-bounds exception? No, because *a[high]* contains the pivot, so *i* will stop advancing when *i* == *high* (if not sooner). There is no such assurance for *j*, though, so the "do { j-- }" loop explicitly tests whether "*j* > *low*" before retreating.

Postscript

The journal "Computing in Science & Engineering" did a poll of experts to make a list of the ten most important and influential algorithms of the twentieth century, and it published a separate article on each of the ten algorithms. Quicksort is one of the ten, and it is surely the simplest algorithm on the list. Quicksort's inventor, Sir C. A. R. "Tony" Hoare, received the ACM Turing Award in 1980 for his work on programming languages, and was conferred the title of Knight Bachelor in March 2000 by Queen Elizabeth II for his contributions to "Computing Science."